

Experiment - 1

Aim:

Overview of different advanced ML and DL techniques and their applications.

Theory:

Machine Learning (ML) and Deep Learning (DL) have rapidly evolved over the past decade, becoming central to solving complex real-world problems across various domains, from healthcare to finance. As the field advances, new techniques and methods have emerged, offering more sophisticated approaches to data analysis and pattern recognition. These advanced methods enable machines to learn from data in increasingly efficient, accurate, and scalable ways. This overview explores some of the most influential and advanced techniques in ML and DL, their benefits, and their applications in diverse industries.

1. Ensemble Learning

Ensemble learning is a powerful technique that combines predictions from multiple models to enhance the overall performance and robustness. The key idea behind ensemble methods is that by leveraging different models or variations of a model, the final prediction can be more accurate and stable than any individual model. The two most popular ensemble learning methods are **Bagging** and **Boosting**.

- **Bagging** (Bootstrap Aggregating) involves training multiple instances of the same learning algorithm on different subsets of the training data, which are generated by random sampling with replacement. The results of these models are then aggregated to form the final prediction. A well-known example of this is the **Random Forest** algorithm, which combines multiple decision trees to improve accuracy and reduce overfitting. Bagging is particularly useful when the base model is prone to high variance, such as decision trees. It is commonly used in areas like image classification, financial predictions, and bioinformatics.

- **Boosting** is a sequential ensemble method where each new model is trained to correct the errors made by previous models. Models are built one after another, and the predictions of each model are weighted based on their accuracy. Boosting techniques like **AdaBoost** and **Gradient Boosting** (e.g., **XGBoost**) are widely used for classification and regression problems. These algorithms have been highly effective in various domains, including fraud detection, medical diagnostics, and stock market forecasting, as they can significantly improve the model's predictive power by focusing on the hardest-to-predict examples.

Both methods aim to reduce bias (Boosting) or variance (Bagging) in the models, making them indispensable in real-world machine learning tasks where high accuracy is critical.

2. Dimensionality Reduction Techniques

Dimensionality reduction is a crucial step in machine learning and deep learning, especially when dealing with high-dimensional data, which can be computationally expensive and lead to overfitting. These techniques aim to reduce the number of features (or dimensions) while preserving the essential information in the data. Among the most popular dimensionality reduction methods are **Principal Component Analysis (PCA)**, **t-Distributed Stochastic Neighbor Embedding (t-SNE)**, and **Autoencoders**.

- **PCA** is a linear technique that identifies the directions (principal components) in which the data varies the most and projects the data along these components to reduce its dimensionality. PCA is widely used in exploratory data analysis, image compression, and preprocessing for other machine learning algorithms. By reducing the complexity of the data, PCA helps to mitigate noise and improve computational efficiency without losing significant information.
- **t-SNE**, on the other hand, is a non-linear technique primarily used for visualizing high-dimensional datasets in a lower-dimensional space (usually 2D or 3D). Unlike PCA, t-SNE focuses on preserving the local structure of the data, making it ideal for clustering or visualizing complex patterns in data, such as word embeddings in NLP or images in computer vision.

- **Autoencoders** are a type of neural network used for unsupervised learning of compressed representations of data. The network consists of an encoder that compresses the data into a lower-dimensional space and a decoder that reconstructs the original data from this compressed representation. Autoencoders have found applications in anomaly detection, image denoising, and data compression, where the goal is to extract meaningful features from the data in an unsupervised manner.

These dimensionality reduction techniques are essential tools in modern machine learning workflows, especially when working with large datasets or when the data contains irrelevant or redundant features.

3. Deep Learning Architectures

Deep learning has revolutionized various fields by allowing models to learn complex patterns directly from raw data. The most commonly used deep learning architectures include **Convolutional Neural Networks (CNNs)**, **Recurrent Neural Networks (RNNs)**, and **Generative Adversarial Networks (GANs)**.

- **CNNs** are designed to process grid-like data, such as images, by applying convolutional filters that capture local patterns (edges, textures, etc.). These networks have been highly successful in computer vision tasks such as image classification, object detection, and facial recognition. CNNs are composed of layers like convolutional layers, pooling layers, and fully connected layers, which work together to extract hierarchical features from raw data. Applications of CNNs extend beyond images to fields like video analysis, medical imaging, and autonomous vehicles, where accurate object recognition and classification are crucial.
- **RNNs** are designed for sequential data, where the output from previous steps influences the current step. This makes RNNs particularly suitable for tasks like time series forecasting, speech recognition, and natural language processing (NLP). However, traditional RNNs suffer from the **vanishing gradient problem**, where they struggle to learn long-term dependencies. **Long Short-Term Memory (LSTM)** networks were introduced to address this problem by introducing memory cells that can remember

information for long periods. LSTMs have been instrumental in improving performance on tasks such as language translation, speech synthesis, and sentiment analysis.

- **GANs** consist of two neural networks—the **generator** and the **discriminator**—that are trained in opposition to each other. The generator creates fake data, while the discriminator tries to distinguish between real and fake data. Over time, the generator improves and can generate highly realistic data. GANs have been used in image generation (e.g., generating realistic images from textual descriptions), video prediction, and even artistic style transfer.

Deep learning architectures, especially CNNs, RNNs, and GANs, are powerful tools in modern AI systems, enabling applications in diverse fields such as healthcare, entertainment, and autonomous systems.

4. Transfer Learning

Transfer learning is a technique in machine learning where a model developed for a particular task is reused as the starting point for a model on a second task. In deep learning, this typically involves using a pre-trained model on a large dataset (such as ImageNet for images or BERT for text) and fine-tuning it on a smaller, domain-specific dataset. This approach has gained popularity because it allows models to leverage knowledge gained from large datasets, even when the new task has limited data.

For example, in image classification, models like **ResNet** or **VGG**, which have been pre-trained on millions of images, can be fine-tuned to classify new objects in a smaller, specific dataset. This significantly reduces the training time and computational resources required, making it ideal for applications where data collection is expensive or time-consuming, such as in medical imaging, where annotated datasets are limited.

Transfer learning is widely used in fields such as NLP and computer vision, where pre-trained models on large corpora or datasets provide a strong foundation for domain-specific tasks like text classification, sentiment analysis, and object recognition.

Conclusion:

Advanced machine learning and deep learning techniques have significantly transformed the landscape of artificial intelligence, enabling applications in diverse fields such as healthcare, finance, robotics, and entertainment. From ensemble methods that combine multiple models to deep learning architectures like CNNs and GANs that can learn complex patterns from raw data, these techniques are shaping the future of AI. As research continues and computational power increases, the applications and effectiveness of these techniques are expected to expand even further.

Experiment - 2

Aim:

Implement a deep neural network from scratch using TensorFlow or PyTorch, gaining hands-on experience in building complex neural architectures.

Theory:

This project implements an Artificial Neural Network (ANN) for binary classification, inspired by the structure and function of the human brain. ANNs consist of interconnected layers of nodes (neurons) that process data by applying weighted transformations and passing results through activation functions, which introduce non-linearity. This ability to capture complex patterns makes ANNs powerful for tasks like classification and regression. The ANN model in this project includes an input layer with 8 features, two hidden layers of 20 neurons each with ReLU activation, and an output layer for the two classes (diabetes or no diabetes). Using Cross Entropy Loss to calculate prediction errors and the Adam optimizer for weight adjustments, the model iteratively minimizes error and improves accuracy over 500 epochs.

Dataset Description:

The **Pima Indians Diabetes Database** (available on Kaggle) contains medical records for women of Pima Indian heritage, with the goal of predicting diabetes onset based on diagnostic variables. This dataset has 768 entries and 8 feature columns, with one target column, "Outcome," indicating diabetes status (1 for positive, 0 for negative).

Attributes (Features):

- **Pregnancies:** Number of times pregnant.
- **Glucose:** Plasma glucose concentration in an oral glucose tolerance test.
- **Blood Pressure:** Diastolic blood pressure (mm Hg).
- **Skin Thickness:** Triceps skin fold thickness (mm).
- **Insulin:** 2-Hour serum insulin (μ U/ml).
- **BMI:** Body mass index ($\text{weight in kg}/(\text{height in m})^2$).

- **Diabetes Pedigree Function:** Likelihood of diabetes based on family history.
- **Age:** Age in years.
- **Outcome:** The target variable (0 = no diabetes, 1 = diabetes).

Code:

▼ Create An ANN Using Pytorch

```

✓ 3s [3] !kaggle datasets download -d uciml/pima-indians-diabetes-database

📄 Dataset URL: https://www.kaggle.com/datasets/uciml/pima-indians-diabetes-database
License(s): CC0-1.0
Downloading pima-indians-diabetes-database.zip to /content
 0% 0.00/8.91k [00:00<?, ?B/s]
100% 8.91k/8.91k [00:00<00:00, 20.6MB/s]

✓ 0s [4] import zipfile
zip_ref = zipfile.ZipFile('/content/pima-indians-diabetes-database.zip', 'r')
zip_ref.extractall('/content')
zip_ref.close()

✓ 1s [5] import pandas as pd

✓ 0s [6] df = pd.read_csv("/content/diabetes.csv")

✓ [7] df.isnull().sum()

✓ [8] df.info()

✓ 2s [9] import seaborn as sns

✓ [10] sns.pairplot(df, hue='Outcome')

✓ 1s [11] from sklearn.model_selection import train_test_split
x = df.drop('Outcome', axis=1)
y = df['Outcome']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=0)

✓ 4s [12] #Libraries
import torch
import torch.nn as nn
import torch.nn.functional as F

✓ 1s [13] # Creating Tensors
x_train = x_train.values
x_test = x_test.values
y_train = y_train.values
y_test = y_test.values

x_train = torch.FloatTensor(x_train)
x_test = torch.FloatTensor(x_test)
y_train = torch.LongTensor(y_train)
y_test = torch.LongTensor(y_test)

```

```
✓ [14] # Creating model
15
class ANN_Model(nn.Module):

    def __init__(self, input_features=8, hidden1=20, hidden2=20, out_features=2):
        super().__init__()
        self.f_connected1 = nn.Linear(input_features, hidden1)
        self.f_connected2 = nn.Linear(hidden1, hidden2)
        self.out = nn.Linear(hidden2, out_features)

    def forward(self, x):
        x = F.relu(self.f_connected1(x))
        x = F.relu(self.f_connected2(x))
        x = self.out(x)
        return x
```

```
✓ [15] ### instantiate my ANN Model
08
torch.manual_seed(20)
model = ANN_Model()
```

```
✓ [16] model.parameters
```

```
✓ [17] # Backward Propagation
58

loss_function = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.01)
```

```
✓ [18] epochs = 500
final_losses = []
for i in range(epochs):
    i = i + 1
    y_pred = model.forward(x_train)
    loss = loss_function(y_pred, y_train)
    final_losses.append(loss)
    if i%10 == 1:
        print("Epoch number: {} and the loss : {}".format(i,loss.item()))

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

```
✓ [19] import matplotlib.pyplot as plt

# Detach the tensor from the computation graph and convert it to a numpy array
final_losses_np = [loss.detach().numpy() for loss in final_losses]

# Plotting
plt.plot(range(epochs), final_losses_np)
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.show()
```



```

✓ [20] # Prediction
0s
    predictions = []
    with torch.no_grad():
        for i, data in enumerate(x_test):
            y_pred = model(data)
            predictions.append(y_pred.argmax().item())

✓ [21] print(predictions)

✓ [22] from sklearn.metrics import confusion_matrix
    cm = confusion_matrix(y_test, predictions)
    cm

✓ [23] plt.figure(figsize=(10,6))
    sns.heatmap(cm, annot=True)
    plt.xlabel('Actual Values')
    plt.ylabel('Predicted Values')

✓ [28] from sklearn.metrics import classification_report, accuracy_score
    print(classification_report(y_test, predictions))
    acc = accuracy_score(y_test, predictions) * 100
    print(acc)


✓ [25] torch.save(model, 'diabetes.pt')
0s

✓ [26] model = torch.load('diabetes.pt')


✓ [ ] model.eval()

```

Output:



[1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1,



	precision	recall	f1-score	support
0	0.80	0.85	0.82	157
1	0.62	0.54	0.58	74
accuracy			0.75	231
macro avg	0.71	0.69	0.70	231
weighted avg	0.74	0.75	0.74	231

Accuracy is: 74.89177489177489

Conclusion:

The neural network model achieves an accuracy of 74.89%, performing better in identifying non-diabetic cases (class `0`) than diabetic cases (class `1`). The model's precision, recall, and F1-score are higher for class `0`, while a higher number of false negatives indicates some missed diabetic cases. The loss curve shows stable convergence around 200 epochs, suggesting effective learning of the training data.

Experiment - 3

Aim:

Utilize pre-trained models and perform transfer learning to solve real-world problems Efficiently.

Theory:

Transfer learning is a machine learning technique where a model developed for one task is reused as the starting point for a model on a different task. It leverages knowledge learned from large, pre-trained models and adapts it to the current problem, which is particularly beneficial when data for the target task is limited. In this experiment, transfer learning enables us to use a powerful pre-trained model (VGG16) on MRI image data to classify brain MRI scans as either "Tumor" or "No Tumor."

Pre Trained Models and VGG16

The VGG16 model is a deep convolutional neural network trained on ImageNet, a large-scale dataset containing millions of labeled images across thousands of categories. The model's architecture, featuring multiple convolutional layers followed by pooling and fully connected layers, has demonstrated remarkable feature extraction capabilities for a wide range of image classification tasks. In transfer learning, VGG16 can serve as a feature extractor, capturing patterns like edges, shapes, and textures which are then used as inputs for the specific classification task of brain tumor detection.

Dataset Description:

The dataset consists of **Brain MRI images** categorized into two classes: "Tumor" and "No Tumor."

- **Tumor:** Images with visible signs of a brain tumor, showing irregular patterns or abnormalities in brain structure.
- **No Tumor:** Images of healthy brains with consistent structures and no signs of tumors.

Each image is resized to **224x224 pixels** to match the VGG16 input requirements. This dataset setup allows the model to learn to distinguish between normal and abnormal brain scans, aiding in efficient tumor detection.

Code:

```
[1] from google.colab import files
    uploaded = files.upload()

    for fn in uploaded.keys():
        print('User uploaded file "{name}" with length {length} bytes'.format(
            name=fn, length=len(uploaded[fn])))

    # Then move kaggle.json into the folder where the API expects to find it.
    !mkdir -p ~/.kaggle/ && mv kaggle.json ~/.kaggle/ && chmod 600 ~/.kaggle/kaggle.json
```

```
✓ [2] !kaggle datasets download -d navoneel/brain-mri-images-for-brain-tumor-detection
```

```
✓ [4] import tensorflow as tf
      from zipfile import ZipFile
      import os, glob
      import cv2
      from tqdm.notebook import tqdm_notebook as tqdm
      import numpy as np
      from sklearn import preprocessing
      from sklearn.model_selection import train_test_split
      from keras.models import Sequential
      from keras.layers import Convolution2D, Dropout, Dense, MaxPooling2D
      from keras.layers import BatchNormalization
      from keras.layers import MaxPooling2D
      from keras.layers import Flatten
      from zipfile import ZipFile
```

```
✓ [5] file_name = "/content/brain-mri-images-for-brain-tumor-detection.zip"
      with ZipFile(file_name, 'r') as zip:
          zip.extractall()
          print('Done')
```

```
✓ [27] os.chdir("/content/yes")
      X = []
      y = []
      for i in tqdm(os.listdir()):
          img = cv2.imread(i)
          img = cv2.resize(img, (224, 224))
          X.append(img)
          y.append((i[0:1]))
          print(i[0:1])
      os.chdir("/content/no")
      for i in tqdm(os.listdir()):
          img = cv2.imread(i)
          img = cv2.resize(img, (224, 224))
          X.append(img)
      for i in range(1, 99):
          y.append('N')
      print(y)
```

```

✓ [7] %matplotlib inline
import matplotlib.pyplot as plt
plt.figure(figsize=(10, 10))
for i in range(4):
    plt.subplot(1, 4, i+1)
    plt.imshow(X[i], cmap="gray")
    plt.axis('off')
plt.show()

✓ [8] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
print ("Shape of an image in X_train: ", X_train[0].shape)
print ("Shape of an image in X_test: ", X_test[0].shape)

✓ ▶ le = preprocessing.LabelEncoder()
y_train = le.fit_transform(y_train)
y_test = le.fit_transform(y_test)
y_train = tf.keras.utils.to_categorical(y_train, num_classes=2)
y_test = tf.keras.utils.to_categorical(y_test, num_classes=2)
y_train = np.array(y_train)
X_train = np.array(X_train)
y_test = np.array(y_test)
X_test = np.array(X_test)
print("X_train Shape: ", X_train.shape)
print("X_test Shape: ", X_test.shape)
print("y_train Shape: ", y_train.shape)
print("y_test Shape: ", y_test.shape)

[11] from keras.applications import vgg16
img_rows, img_cols = 224, 224
vgg = vgg16.VGG16(weights = 'imagenet',
                    include_top = False,
                    input_shape = (img_rows, img_cols, 3))
# Here we freeze the last 4 layers
# Layers are set to trainable as True by default
for layer in vgg.layers:
    layer.trainable = False
# Let's print our layers
for (i,layer) in enumerate(vgg.layers):
    print(str(i) + " " + layer.__class__.__name__, layer.trainable)

✓ 0s [14] def lw(bottom_model, num_classes):
    """creates the top or head of the model that will be
    placed ontop of the bottom layers"""

    top_model = bottom_model.output
    top_model = GlobalAveragePooling2D()(top_model)
    top_model = Dense(1024,activation='relu')(top_model)
    top_model = Dense(1024,activation='relu')(top_model)
    top_model = Dense(512,activation='relu')(top_model)
    top_model = Dense(num_classes,activation='sigmoid')(top_model)
    return top_model

```

```

✓ [15] from keras.models import Sequential
      from keras.layers import Dense, Dropout, Activation, Flatten, GlobalAveragePooling2D
      from keras.layers import Conv2D, MaxPooling2D, ZeroPadding2D

      from keras.models import Model

      num_classes = 2

      FC_Head = lw(vgg, num_classes)

      model = Model(inputs = vgg.input, outputs = FC_Head)

      print(model.summary())

```

```

✓ 0s [16] from tensorflow.keras.models import Model
      model.compile(optimizer='adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])

```

```

✓ [17] history = model.fit(X_train,y_train,
                        epochs=5,
                        validation_data=(X_test,y_test),
                        verbose = 1,
                        initial_epoch=0)

```

```

[18] %matplotlib inline
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(len(acc))
plt.plot(epochs, acc, 'r', label='Training accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend(loc=0)
plt.figure()
plt.show()

```

```

✓ [26] from sklearn.metrics import precision_score, recall_score, f1_score

      # Calculate metrics
      precision = precision_score(y_true_classes, y_pred_classes)
      recall = recall_score(y_true_classes, y_pred_classes)
      f1 = f1_score(y_true_classes, y_pred_classes)

      print(f"Precision: {precision:.2f}")
      print(f"Recall: {recall:.2f}")
      print(f"F1 Score: {f1:.2f}")

```

```

✓ 0s [19] from sklearn.metrics import classification_report, confusion_matrix
      import seaborn as sns

```

```

✓ [20] # Predict classes on the test set
      y_pred = model.predict(X_test)
      y_pred_classes = np.argmax(y_pred, axis=1)
      y_true_classes = np.argmax(y_test, axis=1)

```

```

✓ [21] print("Classification Report:")
      print(classification_report(y_true_classes, y_pred_classes, target_names=['No Tumor', 'Tumor']))

```

```

✓ [22] conf_matrix = confusion_matrix(y_true_classes, y_pred_classes)
      plt.figure(figsize=(6, 5))
      sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=['No Tumor', 'Tumor'], yticklabels=['No Tumor', 'Tumor'])
      plt.xlabel('Predicted Labels')
      plt.ylabel('True Labels')
      plt.title("Confusion Matrix")
      plt.show()


```

```


✓ [25] plt.figure(figsize=(12, 6))
      for i in range(9):
          plt.subplot(3, 3, i + 1)
          plt.imshow(X_test[i])
          plt.axis('off')
          true_label = 'Tumor' if y_true_classes[i] == 1 else 'No Tumor'
          pred_label = 'Tumor' if y_pred_classes[i] == 1 else 'No Tumor'
          plt.title(f"True: {true_label}\nPred: {pred_label}")
      plt.tight_layout()
      plt.show()


```

Output:




<Figure size 640x480 with 0 Axes>


 Precision: 0.86
Recall: 0.94
F1 Score: 0.90

 Classification Report:


	precision	recall	f1-score	support
No Tumor	0.89	0.76	0.82	33
Tumor	0.86	0.94	0.90	51
accuracy			0.87	84
macro avg	0.88	0.85	0.86	84
weighted avg	0.87	0.87	0.87	84




True: No Tumor
Pred: Tumor




True: Tumor
Pred: Tumor




True: Tumor
Pred: Tumor




True: No Tumor
Pred: No Tumor




True: Tumor
Pred: Tumor




True: No Tumor
Pred: No Tumor




True: No Tumor
Pred: No Tumor



True: No Tumor
Pred: Tumor



True: Tumor
Pred: Tumor



Conclusion:

The model effectively classifies brain MRI images into "Tumor" and "No Tumor" categories with an accuracy of 87%. It demonstrates high precision (0.86) and recall (0.94) for tumor detection, making it reliable for identifying cases with tumors, which is critical in medical diagnostics. The confusion matrix shows some misclassifications, particularly in identifying "No Tumor" cases, but overall, the model performs well, especially in minimizing the risk of missing actual tumors. This balance of accuracy and reliability makes it suitable for real-world medical applications.

Experiment-4

Aim:

Implement GAN to Generate Synthetic Data and Explore Its Application in Image Generation and Data Augmentation.

Theory:

The implementation of Generative Adversarial Networks (GANs) involves leveraging their ability to generate realistic synthetic data that mimics the structure of a given dataset. GANs consist of two neural networks, a Generator and a Discriminator, which are trained in an adversarial manner. The generator aims to create data that resembles the training data, while the discriminator attempts to distinguish between real and synthetic data. This adversarial setup pushes both networks to improve, resulting in high-quality synthetic data generation.

In the case of image datasets like MNIST, the goal is to generate synthetic images of handwritten digits (0-9) that are indistinguishable from real samples. This synthetic data has applications in:

Image Generation: Enhancing datasets with more examples of certain classes, improving model robustness.

Data Augmentation: Expanding the training dataset by adding variability, addressing class imbalance, and improving the performance of machine learning models.

Key Steps in Implementing GANs

1. Data Preprocessing

- Load the MNIST dataset.
- Normalize pixel values to the range $[0, 1]$ for stable training.

2. Define the GAN Architecture:

- **Generator:** A neural network that transforms random noise (latent vector) into synthetic images. The network uses techniques like transpose convolution and activation functions (e.g., ReLU, tanh) to generate realistic images.
- **Discriminator:** A classifier neural network that distinguishes between real and synthetic images. It uses convolutional layers and outputs probabilities through a sigmoid activation function.

3. Training Procedure:

- **Adversarial Training:** Train the generator to produce images that deceive the discriminator and train the discriminator to correctly classify images as real or fake.
- **Loss Functions:**
 - **Generator loss:** Measures how effectively the generator can "fool" the discriminator.
 - **Discriminator loss:** Measures how well the discriminator can distinguish between real and synthetic images.
- Use optimizers like Adam with a carefully tuned learning rate to ensure stability.

4. Evaluate Generated Images:

- Visualize generated samples to assess the generator's performance.
- Use metrics like the Inception Score (IS) or Fréchet Inception Distance (FID) to evaluate the quality and diversity of the synthetic data.

Applications of GAN-Generated Synthetic Data

1. Image Generation:

- Creating new samples to enrich datasets or support artistic/creative projects.

2. Data Augmentation:

- Balancing datasets by generating more data for underrepresented classes.
- Supporting training models with additional synthetic data to improve generalization.

3. Simulation and Research:

- Experimenting with "what-if" scenarios by generating data variations that don't exist in the real dataset.

Dataset Description

The MNIST dataset is a collection of grayscale images of handwritten digits from 0 to 9, commonly used in image processing and machine learning.

- **Image Size:** 28x28 pixels.
- **Channels:** 1 (grayscale).
- **Number of Classes:** 10 (digits 0-9).
- **Dataset Split:**
 - Training set: 60,000 images.
 - Test set: 10,000 images.

Code:

```
import tensorflow as tf
from tensorflow.keras.layers import (Dense, BatchNormalization, LeakyReLU, Reshape, Conv2DTranspose, Conv2D, Dropout, Flatten)
import matplotlib.pyplot as plt

(train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()

train_images = train_images.reshape(train_images.shape[0], 28, 28, 1).astype('float32')
test_images = test_images.reshape(test_images.shape[0], 28, 28, 1).astype('float32')

train_images = (train_images - 127.5) / 127.5
test_images = (test_images - 127.5) / 127.5

train_labels = train_labels.astype('float32')
test_labels = test_labels.astype('float32')

Buffer_size = 60000
Batch_size = 256

train_dataset = tf.data.Dataset.from_tensor_slices(train_images).shuffle(Buffer_size).batch(Batch_size)

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 2s 0us/step
```

```

def generator():
    model = tf.keras.Sequential()
    model.add(Dense(7*7*256, use_bias=False, input_shape=(100,)))
    model.add(BatchNormalization())
    model.add(LeakyReLU())

    model.add(Reshape((7, 7, 256)))
    assert model.output_shape == (None, 7, 7, 256)


    model.add(Conv2DTranspose(128, (5, 5), strides=(1, 1), padding='same', use_bias=False))
    assert model.output_shape == (None, 7, 7, 128)
    model.add(BatchNormalization())
    model.add(LeakyReLU())

    model.add(Conv2DTranspose(64, (5, 5), strides=(2, 2), padding='same', use_bias=False))
    assert model.output_shape == (None, 14, 14, 64)
    model.add(BatchNormalization())
    model.add(LeakyReLU())

    model.add(Conv2DTranspose(1, (5, 5), strides=(2, 2), padding='same', use_bias=False, activation='tanh'))
    assert model.output_shape == (None, 28, 28, 1)

    return model

```

 generator = generator()

Run cell (Ctrl+Enter)
cell has not been executed in this session




Model: "sequential"


Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 12544)	1254400
batch_normalization (Batch Normalization)	(None, 12544)	50176
leaky_re_lu (LeakyReLU)	(None, 12544)	0
reshape (Reshape)	(None, 7, 7, 256)	0
conv2d_transpose (Conv2DTranspose)	(None, 7, 7, 128)	819200
batch_normalization_1 (Batch Normalization)	(None, 7, 7, 128)	512
leaky_re_lu_1 (LeakyReLU)	(None, 7, 7, 128)	0
conv2d_transpose_1 (Conv2DTranspose)	(None, 14, 14, 64)	204800
batch_normalization_2 (Batch Normalization)	(None, 14, 14, 64)	256
leaky_re_lu_2 (LeakyReLU)	(None, 14, 14, 64)	0
conv2d_transpose_2 (Conv2DTranspose)	(None, 28, 28, 1)	1600

=====

Total params: 2330944 (8.89 MB)
Trainable params: 2305472 (8.79 MB)
Non-trainable params: 25472 (99.50 KB)

```
noise = tf.random.normal([1,100])
generated_image = generator(noise, training=False)
plt.imshow(generated_image[0, :, :, 0], cmap='gray')
```

 <matplotlib.image.AxesImage at 0x7c82c05a1660>



```
[ ] def discriminator():
    model = tf.keras.Sequential()
    model.add(Conv2D(64, (5, 5), strides=(2, 2), padding='same', input_shape=[28, 28, 1]))
    model.add(LeakyReLU())
    model.add(Dropout(0.3))
    model.add(Conv2D(128, (5, 5), strides=(2, 2), padding='same'))
    model.add(LeakyReLU())
    model.add(Dropout(0.3))
    model.add(Flatten())
    model.add(Dense(1))

    return model
```

```
[ ] discriminator = discriminator()
```

```
discriminator.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 14, 14, 64)	1664
leaky_re_lu_3 (LeakyReLU)	(None, 14, 14, 64)	0
dropout (Dropout)	(None, 14, 14, 64)	0
conv2d_1 (Conv2D)	(None, 7, 7, 128)	204928
leaky_re_lu_4 (LeakyReLU)	(None, 7, 7, 128)	0
dropout_1 (Dropout)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 1)	6273

=====
Total params: 212865 (831.50 KB)
Trainable params: 212865 (831.50 KB)
Non-trainable params: 0 (0.00 Byte)
=====

```
[ ] cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)

def d_loss(real_output, fake_output):
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)
    total_loss = real_loss + fake_loss
    return total_loss

def g_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output), fake_output)

g_optimizer = tf.keras.optimizers.Adam(1e-4)
d_optimizer = tf.keras.optimizers.Adam(1e-4)
```

```
import os

checkpoint_dir = "./training_checkpoints"
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")
checkpoint = tf.train.Checkpoint(generator_optimizer = g_optimizer,
                                  discriminator_optimizer = d_optimizer,
                                  generator = generator,
                                  discriminator = discriminator)
```

```
[ ] EPOCHS = 40
    noise_dim = 100
    num_examples_to_generate = 16

    seed = tf.random.normal([num_examples_to_generate, noise_dim])
```

```
[ ] @tf.function
def train_step(images):

    #creating random noise to feed model
    noise = tf.random.normal([Batch_size, noise_dim])

    # generate images and calculate loss values
    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)
        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)
        gen_loss = g_loss(fake_output)
        disc_loss = d_loss(real_output, fake_output)

    # calculate gradients using loss and model variables
    gradients_of_generator = gen_tape.gradient(gen_loss, generator.trainable_variables)
    gradients_of_discriminator = disc_tape.gradient(disc_loss, discriminator.trainable_variables)

    # process gradients and run optimizer
    g_optimizer.apply_gradients(zip(gradients_of_generator, generator.trainable_variables))
    d_optimizer.apply_gradients(zip(gradients_of_discriminator, discriminator.trainable_variables))
```

```
▶ def generate_and_save_images(model, epoch, test_input):
    predictions = model(test_input, training=False)
    fig = plt.figure(figsize=(4,4))
    for i in range(predictions.shape[0]):
        plt.subplot(4, 4, i+1)
        plt.imshow(predictions[i, :, :, 0] * 127.5 + 127.5, cmap='gray')
        plt.axis('off')

    plt.savefig('image_at_epoch_{:04d}.png'.format(epoch))
    plt.show()
```

```
[ ] import time
from IPython import display

def train(dataset, epochs):
    for epoch in range(epochs):
        start = time.time()

        for image_batch in dataset:
            train_step(image_batch)

        display.clear_output(wait=True)
        generate_and_save_images(generator, epoch + 1, seed)

        if (epoch + 1) % 5 == 0:
            checkpoint.save(file_prefix = checkpoint_prefix)

        print ('Time for epoch {} is {} sec'.format(epoch + 1, time.time()-start))

    display.clear_output(wait=True)
    generate_and_save_images(generator, epochs, seed)
```

```
▶ train(train_dataset, EPOCHS)
```

```
▶ checkpoint.restore(tf.train.latest_checkpoint(checkpoint_dir))
↔ <tensorflow.python.checkpoint.checkpoint.CheckpointLoadStatus at 0x7c82c04f6290>

[ ] import PIL

def display_image(epoch_no):
    return PIL.Image.open('image_at_epoch_{:04d}.png'.format(epoch_no))
display_image(EPOCHS)
```



```

import glob
import imageio

anim_file = 'dcgan.gif'

with imageio.get_writer(anim_file, mode='I') as writer:
    filenames = glob.glob('image*.png')
    filenames = sorted(filenames)
    for filename in filenames:
        image = imageio.imread(filename)
        writer.append_data(image)

display.Image(open(anim_file, 'rb').read())

```

Output:



Conclusion:

In this practical, we implemented a GAN using the MNIST dataset to generate synthetic numeral images, successfully creating realistic digits that expanded our dataset. This experiment highlighted GANs' effectiveness for data augmentation, especially in enhancing model training with additional synthetic data.

Experiment-5

Aim:

Apply NLP techniques to process and analyze textual data, including sentiment analysis & named entity recognition.

Theory:

The analysis of audit comments involves several key steps to extract meaningful insights. First, data preprocessing is crucial for preparing the dataset. This step involves loading the dataset and cleaning the text by removing unnecessary characters, converting the text to lowercase, and eliminating common stop words that do not contribute meaningfully to sentiment or entity recognition. Following this, sentiment analysis is applied to classify each comment as positive, negative, or neutral, giving an initial understanding of the overall tone of the comments. This can be achieved using pre-trained models like those from Hugging Face or tools such as NLTK's VADER, which effectively label sentiments in financial and textual data. Next, Named Entity Recognition (NER) identifies and extracts organization names mentioned within the comments, allowing us to associate sentiments directly with specific entities. Using pre-trained NER models, we can tag and isolate these organizations. In the final analysis phase, the sentiment distribution for each organization is calculated, highlighting organizations with predominantly positive, negative, or neutral sentiments. These results are visualized using charts, providing a clear overview of sentiment trends across organizations. This approach not only aids in filtering entities based on sentiment but also offers valuable insights into patterns within the audit comments.

Sentiment Analysis focuses on determining the emotional tone of each audit comment. By using a pre-trained sentiment model, we classify each comment as positive, negative, or neutral. This allows us to quantify the sentiment across all comments and identify trends in the tone of feedback given by auditors. The sentiment scores help in pre-filtering organizations, highlighting those that might need further analysis due to high positive or negative feedback.

Named Entity Recognition (NER), on the other hand, is used to identify and extract organization names mentioned within the audit comments. By applying a pre-trained NER model, we can tag these entities, linking each comment to specific organizations. This enables us to calculate sentiment distributions for each organization, helping to pinpoint entities with consistently high positive or negative sentiments, thus assisting in trend identification and deeper analysis.

Dataset Description: Audit Comments Dataset -

This dataset, sourced from Hugging Face, contains textual comments from auditors with associated sentiment labels. The dataset is used to analyze the emotional tone of audit feedback toward various organizations. Below is a description of each field in the dataset:

- **Comment Text:** Contains the text of each audit comment, where auditors provide feedback or observations. This field is the primary text data used for both sentiment analysis and named entity recognition (NER) to identify organizations mentioned and assess sentiment trends.
- **Label:** A numerical sentiment label assigned to each comment, categorizing the sentiment as follows:
 - **2** – Positive: The comment expresses a positive sentiment toward the organization.
 - **1** – Neutral: The comment reflects a neutral tone without a strong positive or negative opinion.
 - **0** – Negative: The comment indicates a negative sentiment, often highlighting concerns or issues.


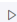
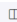
This dataset structure allows for effective sentiment analysis and NER to extract actionable insights on auditor feedback trends across organizations.

Code:

```
sentiment-and-ner-analysis-of-audit-comments.ipynb •    ...  
sentiment-and-ner-analysis-of-audit-comments.ipynb > ...  
+ Code + Markdown | ▶ Run All ⏹ Restart 🧹 Clear All Outputs 📄 Variables 📖 Outline ... base (Python 3.11.7)  
  
[2]   
... # Import the Libraries  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
from IPython.display import HTML  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score  
from IPython.display import display  
  
# tensorflow tensors  
import tensorflow as tf  
  
... /opt/conda/lib/python3.10/site-packages/scipy/_lib/_util.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected  
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

Load the dataset from Hugging Face

```
[3]   
... from datasets import load_dataset  
  
# import the auditor comment dataset from Hugging Face  
dataset = load_dataset("FinanceInc/auditor_sentiment")  
  
# label: a label corresponding to the class as a string:  
# 2 - positive; 1 - neutral; 0 - negative'  
  
... Downloading: 0%| | 0.00/800 [00:00<?, ?B/s]
```


```
sentiment-and-ner-analysis-of-audit-comments.ipynb •    ...  
sentiment-and-ner-analysis-of-audit-comments.ipynb > ...  
+ Code + Markdown | ▶ Run All ⏹ Restart 🧹 Clear All Outputs 📄 Variables 📖 Outline ... base (Python 3.11.7)  
  
... Downloading data: 0%| | 0.00/327k [00:00<?, ?B/s]  
... Extracting data files: 0%| | 0/2 [00:00<?, ?it/s]  
... Dataset parquet downloaded and prepared to /root/.cache/huggingface/datasets/parquet/demo-org--auditor_review-deda15e5348bd334/0.0.0/0b6d5799bb72eb24ad7fc7be72l  
... 0%| | 0/2 [00:00<?, ?it/s]  
  
[4]   
... # Train dataset  
dataset["train"]  
  
... Dataset({  
  features: ['sentence', 'label'],  
  num_rows: 3877  
})  
  
[5]   
... # Test dataset  
dataset["test"]  
  
... Dataset({  
  features: ['sentence', 'label'],  
  num_rows: 969  
})
```

Distribution of Sentiment Labels

Distribution of Sentiment Labels

```
# Distribution of Target label
label_list = list(dataset['train']['label'])
df = pd.DataFrame(label_list, columns=['label'])

# to display the bar charts by target values
rst = df['label'].value_counts().plot(kind='bar')
```



```
# List out some positive comments
positive_dataset[0:5]
```

[0] ('sentence': ['Altia 's operating profit jumped to EUR 47 million from EUR 6.6 million .',
'The agreement was signed with Biohit Healthcare Ltd , the UK-based subsidiary of Biohit Oy] , a Finnish public company which develops , manufactures and marl
'kesko pursues a strategy of healthy , focused growth concentrating on sales and services to consumer-customers .',
'Elcoteq 's stock of orders has stabilised in the past weeks , Mr Krippel said .',
'UPM-Kymmene has generated seventeen consecutive quarters of positive Cash Flow from Operations .'],
'label': [2, 2, 2, 2, 2])

```
# List out some negative comments
negative_dataset[0:5]
```

[0] ('sentence': ['Operating loss totalled EUR 0.9 mn , down from a profit of EUR 2.7 mn .',
'Several large stocks tacked lower , however .',
'As a result some 20 persons will no longer be needed .',
'In addition the production personnel of the Sport Division have been given a temporary lay-off warning .',
'Konecranes has previously communicated an estimated reduction of about 1,600 employees on group level in 2009 .'],
'label': [0, 0, 0, 0, 0])

```
# Parameters
seq_len = 512
batch_size = 12

num_samples = len(dataset['train'])

# to prepare the array for target value
print('Training array of size {}'.format(num_samples, seq_len))
```

[10] Training array of size (3877 x 512)


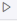
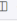
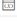

Load the Transformer Model

```
# Checkpoint for the Transformer: "Transformer": Unknown word.
check_point = 'bert-base-uncased'

# Import the Tokenizer of Bert
from transformers import BertTokenizer

# Load the pre-trained model
tokenizer = BertTokenizer.from_pretrained(check_point)
```

[11]

```
sentiment-and-ner-analysis-of-audit-comments.ipynb •    ...
sentiment-and-ner-analysis-of-audit-comments.ipynb > ...
+ Code + Markdown | ▶ Run All ⏹ Restart ⚙ Clear All Outputs  Variables  Outline ...
base (Python 3.11.7)

# tokenize the audit comments in the train dataset
tokens = tokenizer(dataset['train']['sentence'], max_length=seq_len, truncation=True,
padding='max_length', add_special_tokens=True,
return_tensors='pt')

[12] Python

# Input ids
tokens['input_ids']

[13] Python

... tensor([[ 101, 12456, 2401, ..., 0, 0, 0],
[ 101, 1996, 3820, ..., 0, 0, 0],
[ 101, 17710, 21590, ..., 0, 0, 0],
...,
[ 101, 1996, 3189, ..., 0, 0, 0],
[ 101, 24797, 1011, ..., 0, 0, 0],
[ 101, 1996, 2974, ..., 0, 0, 0]])

# Attention masks
tokens['attention_mask']

[14] Python

... tensor([[1, 1, 1, ..., 0, 0, 0],
[1, 1, 1, ..., 0, 0, 0],
[1, 1, 1, ..., 0, 0, 0],
...,
[1, 1, 1, ..., 0, 0, 0],
[1, 1, 1, ..., 0, 0, 0],
[1, 1, 1, ..., 0, 0, 0]])
```

4. Preparation of Target Labels

```
sentiment-and-ner-analysis-of-audit-comments.ipynb •    ...
sentiment-and-ner-analysis-of-audit-comments.ipynb > ...
+ Code + Markdown | ▶ Run All ⏹ Restart ⚙ Clear All Outputs  Variables  Outline ...
base (Python 3.11.7)

# first extract the list of sentiment values, i.e. 0 - 2
arr = pd.DataFrame(dataset['train']['label'], columns=['label'])
arr_values = arr['label'].sort_values().unique().tolist()
print('No. of distinct label values : {}'.format(arr_values))

[15] Python

... No. of distinct label values : [0, 1, 2]

# Initialize the zero array based on the size of the input_ids / attention masks
labels = np.zeros((num_samples, len(arr_values)))
print('Size of the target label : {}'.format(labels.shape))

[16] Python

... Size of the target label : (3877, 3)

# Assign the sentiment values "sentimen": Unknown word.
labels[np.arange(num_samples), arr['label'].tolist()] = 1 "arange": Unknown word.
labels

[17] Python

... array([[0., 0., 1.],
[0., 0., 1.],
[0., 0., 1.],
...,
[0., 0., 1.],
[0., 1., 0.],
[0., 1., 0.]])
```

```
sentiment-and-ner-analysis-of-audit-comments.ipynb
+ Code + Markdown | ▶ Run All ⏹ Restart 🧹 Clear All Outputs | 📄 Variables 📖 Outline ...
base (Python 3.11.7)

5. Sentiment Model Training

# Create the tensorflow tensors from Input Ids, Attention Masks and Labels
ds = tf.data.Dataset.from_tensor_slices((tokens['input_ids'], tokens['attention_mask'], labels))
ds.take(1)

[18] <_TakeDataset element_spec=(TensorSpec(shape=(512,), dtype=tf.int64, name=None), TensorSpec(shape=(512,), dtype=tf.int64, name=None), TensorSpec(shape=(3,), dt

# Conver the tensors from 3d array of (3) tuples into 2d array of (2,1) tuples "Conver": Unknown word.
def map_func(input_ids, masks, labels):
    return {'input_ids': input_ids, 'attention_mask': masks, labels

# then we use the dataset map method to apply this transformation
ds = ds.map(map_func)
# print(len(dataset))

ds = ds.shuffle(10000).batch(batch_size, drop_remainder=True)
# print(len(dataset))

ds.take(1)

[19] <_TakeDataset element_spec=({'input_ids': TensorSpec(shape=(12, 512), dtype=tf.int64, name=None), 'attention_mask': TensorSpec(shape=(12, 512), dtype=tf.int64,

# Create the train and test datasets for Transformer model training
split = 0.8

[20]
```

```
sentiment-and-ner-analysis-of-audit-comments.ipynb
+ Code + Markdown | ▶ Run All ⏹ Restart 🧹 Clear All Outputs | 📄 Variables 📖 Outline ...
base (Python 3.11.7)

Build and Train the Sentiment Model

# Load the Transformer model
from transformers import TFAutoModel

# Set up the pre-trained model of Bert "tained": Unknown word.
model = TFAutoModel.from_pretrained(check_point)

[24] Could not render content for 'application/vnd.jupyter.widget-view+json'
['model_id':'e96e5f90af94203ac1e9703b5fc5faf', 'version_major':2, 'version_minor':0]

Some weights of the PyTorch model were not used when initializing the TF 2.0 model TFBertModel: ['cls.predictions.transform.dense.weight', 'cls.predictions.tran
- This IS expected if you are initializing TFBertModel from a PyTorch model trained on another task or with another architecture (e.g. initializing a TFBertFor
- This IS NOT expected if you are initializing TFBertModel from a PyTorch model that you expect to be exactly identical (e.g. initializing a TFBertForSequenceC
All the weights of TFBertModel were initialized from the PyTorch model.
If your task is similar to the task the model of the checkpoint was trained on, you can already use TFBertModel for predictions without further training.

# we can view the pre-trained Bert model using the summary method "tained": Unknown word.
model.summary()

[25] Model: "tf_bert_model"

Layer (type) Output Shape Param #
=====
bert (TFBertMainLayer) multiple 109482240
=====
Total params: 109,482,240
```

Model Architecture and Structure

```
[26] # Two more input layers - input IDs and Attention Masks
input_ids = tf.keras.layers.Input(shape=(512,), name='input_ids', dtype='int32')
mask = tf.keras.layers.Input(shape=(512,), name='attention_mask', dtype='int32')

# Embedding layer of Bert model
embeddings = model.bert(input_ids, attention_mask=mask)[1] # access final activations (already max-pooled) [1]
# convert bert embeddings into 5 output classes
layer_1 = tf.keras.layers.Dense(1024, activation='relu')(embeddings)
# Dropout layer with a rate of 0.5
x = tf.keras.layers.Dropout(0.1)(layer_1)
layer_2 = tf.keras.layers.Dense(3, activation='softmax', name='outputs')(layer_1) # the sentiment has 3 labels

# initialize model with input layers of Input IDs and Attention Masks and output layer of Outputs
sentiment_model = tf.keras.Model(inputs=[input_ids, mask], outputs=layer_2)

# (optional) freeze bert layer - to decide if the embedding layer is re-trained.
sentiment_model.layers[2].trainable = True

# print the updated model summary
sentiment_model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_ids (InputLayer)	[(None, 512)]	0	[]

Model Fine-tuning with Additional Layers

```
[28] # Training parameters
optimizer = tf.keras.optimizers.Adam(learning_rate=1e-5, decay=1e-6)
loss = tf.keras.losses.CategoricalCrossentropy()
acc = tf.keras.metrics.CategoricalAccuracy('accuracy')

sentiment_model.compile(optimizer=optimizer, loss=loss, metrics=[acc])

# Re-train the model
history = sentiment_model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=3
)
```

Epoch 1/3
258/258 [=====] - 237s 879ms/step - loss: 0.6296 - accuracy: 0.7251 - val_loss: 0.3359 - val_accuracy: 0.8615
Epoch 2/3
258/258 [=====] - 225s 872ms/step - loss: 0.3332 - accuracy: 0.8640 - val_loss: 0.2080 - val_accuracy: 0.9244
Epoch 3/3
258/258 [=====] - 225s 872ms/step - loss: 0.2268 - accuracy: 0.9170 - val_loss: 0.1769 - val_accuracy: 0.9449

```
[30] # free up memory
del train_ds
del val_ds
```


6. Sentiment Analysis

```
> # The test data of one instance only
def prep_data(text):
    tokens = tokenizer.encode_plus(text, max_length=512,
                                   truncation=True, padding='max_length',
                                   add_special_tokens=True, return_token_type_ids=False,
                                   return_tensors='tf')
    # tokenizer returns int32 tensors, we need to return float64, so we use tf.cast
    return {'input_ids': tf.cast(tokens['input_ids'], tf.float64),
            'attention_mask': tf.cast(tokens['attention_mask'], tf.float64)}
```

```
[31] Python
```

```
# The test data in array
def prep_arr_data(text_array):
    tokens = tokenizer.batch_encode_plus(text_array, max_length=512,
                                         truncation=True, padding='max_length',
                                         add_special_tokens=True, return_token_type_ids=False,
                                         return_tensors='tf')
    # tokenizer returns int32 tensors, we need to return float64, so we use tf.cast
    return {'input_ids': tf.cast(tokens['input_ids'], tf.float64),
            'attention_mask': tf.cast(tokens['attention_mask'], tf.float64)}
```

```
[32] Python
```

```
# Create a dictionary to map label indices to label names
label_mapping = {
    0: 'Negative',
    1: 'Neutral',
    2: 'Positive'
}
```

```
[33] Python
```

0 b · span Spaces: 4 Cell 42 of 86 Go Live Quokka ✓ Spell ✓ Prettier {}

Single Test Case

```
> text_array = [
    "XXX's strong financial performance exceeded expectations, driven by robust sales growth and effective cost management.",
    "XXX demonstrated resilience in challenging economic conditions, maintaining profitability through efficient expense management.",
    "XXX's net income declined due to increased operating costs and lower sales volumes. Cost-cutting measures are being implemented.",
    "XXX maintains a strong liquidity position with healthy cash reserves, providing a solid foundation for future investments.",
    "XXX's profitability has steadily increased, attributed to successful product launches and expanded market reach.",
    "XXX's debt management efforts resulted in decreased debt levels and improved debt-to-equity ratio."
]
```

```
probs = sentiment_model.predict(prepare_data(text_array))
pred = np.argmax(probs, axis=1)
```

```
# Sentiment output
for count, prediction in enumerate(pred):
    print('The sentiment for test case {} is {}'.format(count, label_mapping.get(prediction)))
```

```
[34] Python
```

```
1/1 [=====] - 3s 3s/step
The sentiment for test case 0 is Positive.
The sentiment for test case 1 is Positive.
The sentiment for test case 2 is Negative.
The sentiment for test case 3 is Positive.
The sentiment for test case 4 is Positive.
The sentiment for test case 5 is Positive.
```

0 b · span Spaces: 4 Cell 42 of 86 Go Live Quokka ✓ Spell {}

sentiment-and-ner-analysis-of-audit-comments.ipynb • 6. Sentiment Analysis • # The test data of one instance only

Code + Markdown Run All Restart Clear All Outputs Variables Outline ... base (Python 3.11.7)

Dataset of Sentences

```
# Sentiment prediction
probs = sentiment_model.predict(prepare_data(dataset['test']['sentence'])) "probs": Unknown word.

# import numpy as np
pred = np.argmax(probs, axis=1) "argmax": Unknown word.
```

[35] Python

... 31/31 [=====] - 21s 670ms/step

```
df_test = pd.DataFrame(dataset['test']['label'], columns = ['label'])
df_pred = pd.DataFrame(pred, columns = ['label'])
```

[36] Python

+ Code + Markdown

```
# Calculate accuracy
accuracy = accuracy_score(df_pred, df_test) * 100

# Print the accuracy
print(f"Prediction Accuracy: {accuracy:.2f}%")
```

[37] Python


... Prediction Accuracy: 85.55%

Sentiment output

W 0 b · span Spaces: 4 Cell 42 of 86 Go Live Quokka ✓ Spell {}

sentiment-and-ner-analysis-of-audit-comments.ipynb • 6. Sentiment Analysis • # The test data of one instance only

Code + Markdown Run All Restart Clear All Outputs Variables Outline ... base (Python 3.11.7)



Sentiment Category	Count
Negative	125
Neutral	200
Positive	200

```
# Sentiment for positive comments
positive_pred = df_pred[df_pred['label'] == 2][0:5]
# positive_pred

# Sample output
for idx in range(5):
    print('The sentiment is \033[1mPositive\033[0m for sentence of "{}".'.format(dataset['test']['sentence'][positive_pred.index[idx]]))
    print("")
```

[40] Python

... The sentiment is **Positive** for sentence of "Teliasonera TISN said the offer is in line with its strategy to increase its ownership in core business holdings and

The sentiment is **Positive** for sentence of "STORA ENSO , NORSKE SKOG , M-REAL , UPM-KYMMENE Credit Suisse First Boston (CFSB) raised the fair value for shares

The sentiment is **Positive** for sentence of "Clothing retail chain Sephora's sales increased by 8 % to EUR 155,2 mn , and operating profit rose to EUR 31,1 mn

The sentiment is **Positive** for sentence of "Lifetree was founded in 2000 , and its revenues have risen on an average by 40 % with margins in late 30s .",

The sentiment is **Positive** for sentence of "Nordea Group 's operating profit increased in 2010 by 18 percent year-on-year to 3,64 billion euros and total revenue

W 0 b · span Spaces: 4 Cell 42 of 86 Go Live Quokka ✓ Spell {}

```
sentiment-and-ner-analysis-of-audit-comments.ipynb • [Python 3.11.7]
6. Sentiment Analysis > # The test data of one instance only
+ Code + Markdown | ▶ Run All ⌂ Restart 🗑 Clear All Outputs 📄 Variables 📖 Outline ...

# Sample output
for idx in range(5):
    print('The sentiment is \033[1mNegative\033[0m for sentence of "{}.format(dataset['test']['sentence'][negative_pred.index[idx]])
    print('')

[41] Python
... The sentiment is Negative for sentence of "Compared with the FTSE 100 index , which rose 51.5 points ( or 0.9 % ) on the day , this was a relative price change
The sentiment is Negative for sentence of "Calls to the switchboard and directory services have decreased significantly since our employees now have up-to-date
The sentiment is Negative for sentence of "One of the challenges in the oil production in the North Sea is scale formation that can plug pipelines and halt produ
The sentiment is Negative for sentence of "Profit before taxes amounted to EUR 56.5 mn , down from EUR 232.9 mn a year ago .".
The sentiment is Negative for sentence of "Vaisala also said it expects net sales of EUR 253.2 million for 2010 , compared with EUR 252.2 million recorded in 20

[42] Python
# Sentiment for neutral comments
neutral_pred = df_pred[df_pred['label'] == 1][0:5]
# negative_pred

# Sample output
for idx in range(5):
    print('The sentiment is \033[1mNeutral\033[0m for sentence of "{}.format(dataset['test']['sentence'][negative_pred.index[idx]])
    print('')

[42] Python
... The sentiment is Neutral for sentence of "Compared with the FTSE 100 index , which rose 51.5 points ( or 0.9 % ) on the day , this was a relative price change
The sentiment is Neutral for sentence of "Calls to the switchboard and directory services have decreased significantly since our employees now have up-to-date
The sentiment is Neutral for sentence of "One of the challenges in the oil production in the North Sea is scale formation that can plug pipelines and halt produ
The sentiment is Neutral for sentence of "Profit before taxes amounted to EUR 56.5 mn , down from EUR 232.9 mn a year ago .".
The sentiment is Neutral for sentence of "Vaisala also said it expects net sales of EUR 253.2 million for 2010 , compared with EUR 252.2 million recorded in 20
```

```
sentiment-and-ner-analysis-of-audit-comments.ipynb • [Python 3.11.7]
7. Named Entity Recognition (NER)
+ Code + Markdown | ▶ Run All ⌂ Restart 🗑 Clear All Outputs 📄 Variables 📖 Outline ...

... The sentiment is Neutral for sentence of "Compared with the FTSE 100 index , which rose 51.5 points ( or 0.9 % ) on the day , this was a relative price change
The sentiment is Neutral for sentence of "Calls to the switchboard and directory services have decreased significantly since our employees now have up-to-date
The sentiment is Neutral for sentence of "One of the challenges in the oil production in the North Sea is scale formation that can plug pipelines and halt produ
The sentiment is Neutral for sentence of "Profit before taxes amounted to EUR 56.5 mn , down from EUR 232.9 mn a year ago .".
The sentiment is Neutral for sentence of "Vaisala also said it expects net sales of EUR 253.2 million for 2010 , compared with EUR 252.2 million recorded in 20

[43] Python
# Parameter
ner_name = 'test'

[44] Python
# Load the model
from transformers import AutoTokenizer, AutoModelForTokenClassification
from transformers import pipeline

checkpoint_2 = "dslim/bert-base-NER" "dslim": Unknown word.

# Load the pre-trained Bert model for NER
ner_tokenizer = AutoTokenizer.from_pretrained(checkpoint_2)
ner_model = AutoModelForTokenClassification.from_pretrained(checkpoint_2)

[44] Python
b span Spaces: 4 Cell 56 of 85 Go Live Quokka Spell
```

7. Named Entity Recognition (NER)

```
sentiment-and-ner-analysis-of-audit-comments.ipynb • 8. Formation of Organization Entity
sentiment-and-ner-analysis-of-audit-comments.ipynb > M 8. Formation of Organization Entity
+ Code + Markdown | ▶ Run All ⏮ Restart ⚙ Clear All Outputs | 📄 Variables 📄 Outline ... base (Python 3.11.7)

# The Bert model returns B-ORG and I-ORG, the B-ORG is the beginning of an organization entity identified by the model "organization": Unknown word.
# The I-ORG is a linked token to the precedent B-ORG or I-ORG token.
# The concatenation of consecutive B-ORG and subsequent I-ORG will form the name of the organization entity "concatation": Unknown word.

entity_dict = {}
entities = []
current_entity = ""
for idx in range(len(ner_results)):
    for token in ner_results[idx]:
        if token['entity'] == "B-ORG":
            if token['word'][0:2] == "##":
                current_entity += token['word'][2:]
            elif current_entity != '':
                entities.append(current_entity)
                current_entity = token['word']
            else:
                current_entity = token['word']
        elif token['entity'] == "I-ORG":
            if token['word'][0:2] == "##":
                current_entity += token['word'][2:]
            else:
                current_entity += " " + token['word']

    entities.append(current_entity) # Append the last token to the entity list

entity_dict[idx] = entities

# empty the array and current entity
entities = []
```

```
sentiment-and-ner-analysis-of-audit-comments.ipynb • 8. Formation of Organization Entity
sentiment-and-ner-analysis-of-audit-comments.ipynb > M 8. Formation of Organization Entity
+ Code + Markdown | ▶ Run All ⏮ Restart ⚙ Clear All Outputs | 📄 Variables 📄 Outline ... base (Python 3.11.7)

... No. of records in the NER results : 969

# After NER, null return will be removed, and we only need the records with return of organization names

# All records and organization names after the NER
entity_2d_array = [[k, v] for k, values in entity_dict.items() for v in values]

# Remove those records with null return
entity_2d_array_dedup = [sublist for sublist in entity_2d_array if all(item != '' for item in sublist)] "dedup": Unknown word.

print('No. of organizations in the original NER: {}'.format(len(entity_2d_array)))
print('No. of organizations after the de-duplication: {}'.format(len(entity_2d_array_dedup))) "dedup": Unknown word.

[48] Python
... No. of organizations in the original NER: 1366
No. of organizations after the de-duplication: 935

# List some of the NER results
entity_2d_array_dedup[:10] "dedup": Unknown word.

[49] Python
... [[0, 'TeliaSonera TLSN'],
[0, 'Eesti Telekom'],
[1, 'STORA ENSO'],
[1, 'NORSKE SKOG'],
[1, 'M R'],
[1, 'KYMME Credit Suisse First Boston'],
[1, 'CFSB'],
[3, 'Lifetree'],
[4, 'Nordea Group'],
[6, 'Lithuanian Brewers ' Association']]
```

Output -

sentiment-and-ner-analysis-of-audit-comments.ipynb

sentiment-and-ner-analysis-of-audit-comments.ipynb > 9. Analysis of Audit Comments

+ Code + Markdown | ▶ Run All ↺ Restart ≡ Clear All Outputs 📄 Variables ≡ Outline ...

base (Python 3.11.7)




Sentiment and NER Results

df_ner_sentiment = pd.DataFrame(merged_array, columns=['doc_id', 'org_name', 'sentiment_label'])
df_ner_sentiment

[57] Python

	doc_id	org_name	sentiment_label
0	0	TeliaSonera TLSN	2
1	0	Eesti Telekom	2
2	1	STORA ENSO	2
3	1	NORSKE SKOG	2
4	1	M R	2
...
930	955	Wireless Solutions	2
931	957	Danske Bank A	2
932	957	DANKE DC	2
933	964	Ragutis	0
934	964	Olvi	0

935 rows × 3 columns

  ... 

sentiment-and-ner-analysis-of-audit-comments.ipynb

sentiment-and-ner-analysis-of-audit-comments.ipynb > 9. Analysis of Audit Comments

+ Code + Markdown | ▶ Run All ↺ Restart ≡ Clear All Outputs 📄 Variables ≡ Outline ...

base (Python 3.11.7)

filter_by_sentiment (2, 'Positive')

[59] Python

=====

Top 10 Companies with Positive Audit Comments based on Sentiment Analysis

=====

```
org_name
ADP News      4
E              4
Elcoteq       3
Kesko         3
Nokia         3
Aspo          3
M - real      3
KCI Konecranes 3
Talvivaara    3
Outotec       3
Name: doc_id, dtype: int64
```

filter_by_sentiment (0, 'Negative')

[60] Python

=====

Top 10 Companies with Negative Audit Comments based on Sentiment Analysis

=====

```
org_name
Talentum      4
Pretax        3
Scanfil       3
L & T         2
```

0 b · span Spaces: 4 Cell 76 of 84 Go Live Quokka ✓ Spell

```
sentiment-and-ner-analysis-of-audit-comments.ipynb • 
sentiment-and-ner-analysis-of-audit-comments.ipynb > 9. Analysis of Audit Comments
+ Code + Markdown | ▶ Run All ↺ Restart ≡ Clear All Outputs | 📄 Variables 📄 Outline ...

... org_name
Talentum          4
Pretax            3
Scanfil           3
L & T             2
EB                2
Cencorp Corporation 2
OMX Helsinki      2
Nordic Aluminium  1
Nobel Biocare     1
Ramirent          1
Name: doc_id, dtype: int64

▶ filter_by_sentiment(1, 'Neutral')
[61]

... =====
Top 10 Companies with Neutral Audit Comments based on Sentiment Analysis
=====

... org_name
Nokia            10
OMX Helsinki     7
Glaston          6
Nordea           5
Alma Media       5
Vacon            4
Teleste          4
Newstex          3
Aspo             3
Technopolis      3
Name: doc_id, dtype: int64
```

Conclusion:

In conclusion, this project successfully applied sentiment analysis and named entity recognition (NER) to auditor comments, enabling an automated assessment of sentiment trends across various organizations. By identifying positive, negative, and neutral sentiments and linking them to specific entities, we gained valuable insights into how organizations are perceived. This approach streamlines the review process, allowing for quicker identification of organizations with significant positive or negative feedback for deeper analysis.